

Speculative Execution Resilient Cryptography

Rui Fernandes

Supervisors: Bernardo Portela, José Bacelar Almeida, Tiago Oliveira

July 2023

Master's degree in Information Security
Faculty of Sciences of the University of Porto

- **Speculative Execution:** Modern CPUs improve performance by predicting and executing instructions ahead of time;
- **Speculative Execution Attacks:** Exploit speculative execution to execute instructions that would not be executed otherwise. Examples:
 - **Spook.js [1]:** Read sensitive information such as passwords. Affects all Chromium-based browsers;
 - **Foreshadow [2]:** Extract cryptographic keys from Intel SGX enclaves;
 - **Meltdown [3]:** Read kernel memory from the user-space.

Motivation

- We can block speculative execution by inserting LFENCE instructions at specific points in the program;
 - The excessive use of these instructions results in a considerable performance penalty.
- **Alternative:** Speculative Load Hardening (SLH) \Rightarrow harden speculative loads from memory;
- Jasmin is a framework for writing efficient and verified cryptographic software:
 - Cryptographic software is security critical;
 - In Jasmin, we can protect programs against Spectre v1 attacks at the source level by implementing SLH using a type system [4].

Objectives

- Protect `libjbn`, a Jasmin big number library, against speculative execution attacks;
- Extend `libjbn` with a generic implementation of arithmetic operations over elliptic curves.

Background

Side-channel Attacks

Side-channel attacks exploit information leaked through the implementation rather than targeting the underlying cryptographic algorithms:

- **Timing side-channel attacks:** Exploit differences in the execution time of a program:
 - **Constant-Time (CT) Programming:** Control-flow and memory accesses should be independent of secret data;
- **Cache side-channel attacks:** Exploit the fact that accessing data in the cache is faster than accessing data from memory;
- Other side-channels can also be exploited: e.g. power consumption.

Spectre v1 – Bounds Check Bypass

Exploits conditional branch misprediction;

```
if (x < array1_size) { // x is untrusted input  
    index = array1[x] * 4096;  
    y = array2[index];  
}
```

1: Conditional branch misprediction [5]

Speculative Load Hardening

- **Speculative Load Hardening (SLH):** Maintain a predicate indicating whether the execution is misspeculating or not. If it is, this value is then used to “poison” both values and memory addresses of load instructions.
- **Selective Speculative Load Hardening (selSLH):** Not all speculative loads need to be hardened.

Jasmin

Framework for writing cryptographic implementations:

- Verification-friendly;
- Low-level;
- High-assurance and high-speed;
- The compiler is proved to be functionally correct in the Coq proof assistant;
- Automatic safety checker:
 - Memory-Safety;
 - Termination.

Jasmin: Example

```
fn sum(reg ptr u64[100] p) -> reg u64 {  
    reg bool cond;  
    reg u64 sum i;  
  
    sum = 0; i = 0;  
  
    while { cond = (i < 100); } (cond) {  
        sum += p[(int) i];  
        i += 1;  
    }  
  
    return sum;  
}
```

2: Jasmin local function [6]

Control-Flow:

- For loops are fully unrolled:
 - The compiled assembly does not contain any branching instructions;
 - The number of iterations must be known at compile-time;
- While loops/if statements compile to branching instructions.

Functions:

- **Inline functions:** Function calls are replaced by the function body;
- **Local functions:** Compile to `call/ret` instructions;
- **Export functions:** Can be called from other programs – e.g. C programs.

Speculative Type System

Security Levels & Security Types

- Two security levels L and H :
 - L denotes a low security level.
 - H a high security level.
- A security type is a pair of security levels (τ_n, τ_s) :
 - τ_n denotes the security level under normal (i.e. sequential) execution.
 - τ_s denotes the security level of all executions (including misspeculation).
 - (L, L) denotes public data;
 - (H, H) denotes secret data;
 - (L, H) denotes transient data.

Speculative Constant-Time (SCT) & Speculative Type System

- A program is SCT if, for every possible choice of speculation, it does not leak anything beyond what is leaked during sequential execution;
- The type system provides primitives to implement selSLH, ensuring that code is SCT.

Jasmin	Semantics	Compiled to
<code>ms = #init_msf();</code>	<code>ms = 0</code>	<code>lfence; ms = 0;</code>
<code>ms = #update_msf(e, ms);</code>	<code>assert(e)</code>	<code>ms = -1 if !e;</code>
<code>x = #protect(x, ms);</code>	<code>assert(ms == 0)</code>	<code>x = ms;</code>

Table 1: Type system primitives

Speculative Type System – Example

```
fn sum(#msf reg u64 ms, reg ptr u64[100] p)
  -> #msf reg u64, #public reg u64 {
    #public reg bool cond;
    #public reg u64 sum i;

    sum = 0; i = 0;

    while { cond = (i < 100); } (cond) {
      ms = #update_msf(cond, ms);
      sum += p[(int) i];
      i += 1;
    }

    ms = #update_msf(!cond, ms);
    sum = #protect(sum, ms);

    return ms, sum;
  }
```

Speculative Type System – Example

```
fn sum(#msf reg u64 ms, reg ptr u64[100] p)
  -> #msf reg u64, #public reg u64 {
    #public reg bool cond;
    #public reg u64 sum i;

    sum = 0; i = 0;

    while { cond = (i < 100); } (cond) {
      ms = #update_msf(cond, ms);
      sum += p[(int) i];
      i += 1;
    }

    ms = #update_msf(!cond, ms);
    sum = #protect(sum, ms);

    return ms, sum;
}
```

Speculative Type System – Example

```
fn sum(#msf reg u64 ms, reg ptr u64[100] p)
  -> #msf reg u64, #public reg u64 {
    #public reg bool cond;
    #public reg u64 sum i;

    sum = 0; i = 0;

    while { cond = (i < 100); } (cond) {
      ms = #update_msf(cond, ms);
      sum += p[(int) i];
      i += 1;
    }

    ms = #update_msf(!cond, ms);
    sum = #protect(sum, ms);

    return ms, sum;
}
```

Speculative Type System – Example

```
fn sum(#msf reg u64 ms, reg ptr u64[100] p)
  -> #msf reg u64, #public reg u64 {
    #public reg bool cond;
    #public reg u64 sum i;

    sum = 0; i = 0;

    while { cond = (i < 100); } (cond) {
      ms = #update_msf(cond, ms);
      sum += p[(int) i];
      i += 1;
    }

    ms = #update_msf(!cond, ms);
    sum = #protect(sum, ms);

    return ms, sum;
}
```

Implementation

Jasmin big number library;

- Defines a set of basic arithmetic operations for big integers and finite field arithmetic;
- Each number is represented an array of 64-bit unsigned integers;
- Functions are generic on the number of limbs of the numbers;
- The programmer must provide a set of parameters: e.g. number of limbs, size of the finite field.

Source Code Modifications

- If needed, free one register for the misspeculation flag;
- If needed, change function signature to return the misspeculation flag;
- Add security type annotations:
 - Arguments of export functions are transient;
 - Memory addresses are public;
 - Big numbers are secret.
- If needed, protect loads from memory;
- Declassify public values loaded from memory.

Source Code Modifications – Example

```
export fn bn_set0(reg u64 rp) {  
    inline int i;  
  
    for i = 0 to NLIMBS {  
        [rp + 8*i] = 0;  
    }  
}
```

```
export fn bn_set0(#transient reg u64 rp) {  
    _ = #init_msf();  
  
    inline int i;  
  
    for i = 0 to NLIMBS {  
        [rp + 8*i] = 0;  
    }  
}
```

Source Code Modifications – Example

```
// ... implementation omitted for brevity
for i = 0 to NLIMBS {
    t = scalar[(int) i];

    k = 64;
    while { cond = (k > 0); } (cond) {
        ms = #update_msf(cond, ms);
        sk = k; st = t;
        // ... implementation omitted for brevity
        t = st;
        // ... implementation omitted for brevity
        k = sk;
        k = #protect(k, ms);
        k -= 1;
    }
    ms = #update_msf(!cond, ms);
}
```

Performance Evaluation – Number of Cycles (Integer Arithmetic)

Function	CT (Clock Cycles)	SCT (Clock Cycles)	Overhead (%)
bn_addn	60	100	66.67
bn_copy	60	100	66.67
bn_eq	60	100	66.67
bn_muln	100	140	40
bn_set0	60	100	66.67
bn_sqrn	100	140	40
bn_subn	60	100	66.67
bn_test0	60	100	66.67

Table 2: Performance comparison of integer arithmetic functions for 4 limbs

Performance Evaluation – Number of Cycles (Field Arithmetic)

Function	CT (Clock Cycles)	SCT (Clock Cycles)	Overhead (%)
fp_add	80	160	100
fp_expm_noct	35760	35560	-0.56
fp_fromM	120	180	50
fp_inv	63180	63520	0.54
fp_mul	220	280	27.27
fp_sqr	220	280	27.27
fp_sub	80	140	75
fp_toM	140	200	42.86

Table 3: Performance comparison of finite field arithmetic functions for 4 limbs

Performance Evaluation – Number of Cycles (Field Arithmetic)

Function	CT (Clock Cycles)	SCT (Clock Cycles)	Overhead (%)
fp_add	80	160	100
fp_expm_noct	35760	35560	-0.56
fp_fromM	120	180	50
fp_inv	63180	63520	0.54
fp_mul	220	280	27.27
fp_sqr	220	280	27.27
fp_sub	80	140	75
fp_toM	140	200	42.86

Table 3: Performance comparison of finite field arithmetic functions for 4 limbs

Elliptic Curve Cryptography

Elliptic Curve Cryptography: Algebraic Addition

Incomplete addition formulas:

- If $P_1 = \mathcal{O}$, then $P_1 \oplus P_2 = P_2$.
- If $P_2 = \mathcal{O}$, $P_1 \oplus P_2 = P_1$.
- If $P_2 = -P_1$, i.e. if $P_2 = (x_1, -y_1)$, then $P_1 \oplus P_2 = \mathcal{O}$.
- If $P_1 = P_2$, then $P_1 \oplus P_2 = 2 \cdot P_1 = (x_3, y_3)$, where:

$$x_3 = \lambda^2 - 2x_1, \quad y_3 = \lambda(x_1 - x_3) - y_1, \quad \text{where } \lambda = \frac{3x_1^2 + a}{2y_1}$$

- Otherwise, $P_1 \oplus P_2 = P_3 = (x_3, y_3)$, where:

$$x_3 = \lambda^2 - x_1 - x_2, \quad y_3 = \lambda(x_1 - x_3) - y_1, \quad \text{where } \lambda = \frac{y_1 - y_2}{x_1 - x_2}$$

Scalar Multiplication

Given a point $P \in E(\mathbb{F}_p)$ and a scalar $k \in \mathbb{Z}$, we compute the point $Q = k \cdot P$ by repeatedly adding P with itself:

$$Q = k \cdot P = \underbrace{P \oplus P \oplus \dots \oplus P}_{k \text{ terms in the sum}}$$

Scalar Multiplication – Implementation

Require: t -bit scalar $k = (k_{t-1}, \dots, k_1, k_0)_2$, Point $P \in E(\mathbb{F}_p)$

Ensure: $Q = k \cdot P$

```
1:  $Q \leftarrow \mathcal{O}$ 
2: for  $i$  from 0 to  $t - 1$  do
3:   if  $k_i = 1$  then
4:      $Q \leftarrow Q \oplus P$ 
5:   end if
6:    $P \leftarrow 2P$ 
7: end for
8: return  $Q$ 
```

Algorithm 1: Right-to-left binary method for point multiplication

Scalar Multiplication – Implementation

- Vulnerable to side-channel attacks – leaks the binary representation of the scalar k .
- In ECDSA, the public key Q is computed as $Q = k \cdot G$, where k is the private key and G is the base point of the curve.




Figure 1: Power trace of the double-and-add algorithm on a RISC-V CPU [7]

Scalar Multiplication – Montgomery Ladder

Require: t -bit scalar $k = (k_{t-1}, \dots, k_1, k_0)_2$, Point $P \in E(\mathbb{F}_p)$

Ensure: $Q = k \cdot P$

```
1:  $R_0 \leftarrow \mathcal{O}$ 
2:  $R_1 \leftarrow P$ 
3: for  $i$  from 0 to  $t - 1$  do
4:   if  $k_i = 0$  then
5:      $R_1 \leftarrow R_0 \oplus R_1$ 
6:      $R_0 \leftarrow 2 \cdot R_0$ 
7:   else
8:      $R_0 \leftarrow R_0 \oplus R_1$ 
9:      $R_1 \leftarrow 2 \cdot R_1$ 
10:  end if
11: end for
12: return  $R_0$ 
```

Algorithm 2: Montgomery Ladder for point multiplication

Arithmetic Operations Over Elliptic Curves

- Incomplete addition formulas are vulnerable to side-channel attacks and require more effort to protect against Spectre v1 attacks;
- **Complete Addition Formulas:** Compute the sum of *any* two points of $E(\mathbb{F}_p)$:
 - Slower than incomplete addition formulas;
 - Protection against side-channel attacks.

Performance Evaluation – Number of Cycles (Elliptic Curve Arithmetic)

Function	CT (Clock Cycles)	SCT (Clock Cycles)	Overhead (%)
<code>ecc_add</code>	2480	2560	3.23
<code>ecc_branchless_scalar_mul</code>	1732380	1741960	0.55
<code>ecc_double</code>	2200	2280	3.64
<code>ecc_mixed_add</code>	2180	2220	1.83
<code>ecc_normalize</code>	63980	63380	-0.94
<code>ecc_scalar_mul</code>	1178220	1187860	0.82

Table 4: Performance comparison of elliptic curve arithmetic functions for 4 limbs

Performance Evaluation – Number of Cycles (Elliptic Curve Arithmetic)

Function	CT (Clock Cycles)	SCT (Clock Cycles)	Overhead (%)
<code>ecc_add</code>	2480	2560	3.23
<code>ecc_branchless_scalar_mul</code>	1732380	1741960	0.55
<code>ecc_double</code>	2200	2280	3.64
<code>ecc_mixed_add</code>	2180	2220	1.83
<code>ecc_normalize</code>	63980	63380	-0.94
<code>ecc_scalar_mul</code>	1178220	1187860	0.82

Table 4: Performance comparison of elliptic curve arithmetic functions for 4 limbs

Performance Evaluation – Number of Cycles

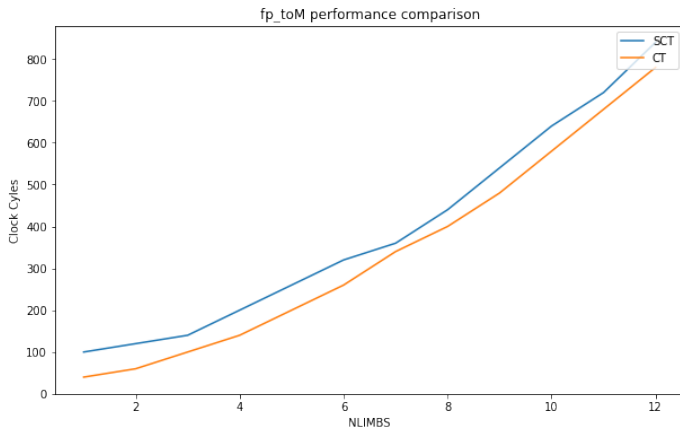


Figure 2: Cycle Count in terms of the number of limbs for the fp_toM function

Conclusion & Future Work

It is possible to protect cryptographic implementations against Spectre v1 with minimal overhead.

Limitations

- The Jasmin compiler is not proved to preserve the protections enforced by the type system \Rightarrow we do not have the guarantee that the compiled assembly is SCT.
- Other side-channels can still be exploited (e.g. Differential Power Analysis attacks).

- Optimized implementations leveraging AVX/AVX2 instructions.
- Formal proofs of correctness in EasyCrypt.

Questions?

- [1] A. Agarwal, S. O'Connell, J. Kim, S. Yehezkel, D. Genkin, E. Ronen, and Y. Yarom, "Spook.js: Attacking Chrome Strict Site Isolation via Speculative Execution," in *2022 IEEE Symposium on Security and Privacy (SP)*, 2022, pp. 699–715.
- [2] J. Van Bulck, M. Minkin, O. Weisse, D. Genkin, B. Kasikci, F. Piessens, M. Silberstein, T. F. Wenisch, Y. Yarom, and R. Strackx, "Foreshadow: Extracting the keys to the Intel SGX kingdom with transient out-of-order execution," in *Proceedings of the 27th USENIX Security Symposium*. USENIX Association, August 2018.
- [3] M. Lipp, M. Schwarz, D. Gruss, T. Prescher, W. Haas, A. Fogh, J. Horn, S. Mangard, P. Kocher, D. Genkin, Y. Yarom, and M. Hamburg, "Meltdown: Reading Kernel Memory from User Space," in *27th USENIX Security Symposium (USENIX Security 18)*, 2018.

- [4] B. A. Shivakumar, G. Barthe, B. Grégoire, V. Laporte, T. Oliveira, S. Priya, P. Schwabe, and L. Tabary-Maujean, "Typing high-speed cryptography against spectre v1," Cryptology ePrint Archive, Paper 2022/1270, 2022, <https://eprint.iacr.org/2022/1270>. [Online]. Available: <https://eprint.iacr.org/2022/1270>
- [5] P. Kocher, J. Horn, A. Fogh, D. Genkin, D. Gruss, W. Haas, M. Hamburg, M. Lipp, S. Mangard, T. Prescher, M. Schwarz, and Y. Yarom, "Spectre Attacks: Exploiting Speculative Execution," in *2019 IEEE Symposium on Security and Privacy (S&P)*, 2019, pp. 1–19.
- [6] "jasmin-lang/jasmin: Language for high-assurance and high-speed cryptography," <https://github.com/jasmin-lang/jasmin>, (Accessed on 14/02/2023).
- [7] U. Banerjee, "Efficient Algorithms, Protocols and Hardware Architectures for Next-Generation Cryptography in Embedded Systems," Thesis, Massachusetts Institute of Technology, Jun. 2021. [Online]. Available: <https://dspace.mit.edu/handle/1721.1/139330>